

RSA Cryptography Optimization for Embedded Systems

Final Project – SENG440, University of Victoria

Waltvin Lee
waltvin@uvic.ca

10th August 2021

Presentation Outline

- Project Specifications
- RSA Public-key Cryptosystem
- Optimizations Performed & Specific Details
- Optimizations that Reduced Performance
- Cycle-count results
- Hardware-assisted solution analysis
- Conclusion
- Invitation for Questions

Project Specifications

Objectives

- Implement RSA Encryption and Decryption targeted at 95-bit keys
- Optimize Montgomery Modular Multiplication
- Optimize 96-bit arithmetic operations
- Analyse hardware-assisted solutions

Platform

- QEMU emulating Fedora VM with ARMv7 Processor
- gcc compiler in the VM

(Rivest–Shamir–Adleman) Public-key Cryptosystem

- Public key pair (E, PQ) published publicly
- Private key D kept secret
- Keys typically are most typically 1024 to 4096 bits

To encrypt

$$C = T^E \bmod PQ, \text{ where } C = \text{ciphertext}, T = \text{plaintext}$$

To decrypt

$$T = C^D \bmod PQ, \text{ where } C = \text{ciphertext}, T = \text{plaintext}.$$

(Rivest–Shamir–Adleman) Public-key Cryptosystem

- Techniques to perform Modular Exponential efficiently
 - Multiply and Square Algorithm
 - Montgomery Modular Multiplication (MMM)
- Modular Exponential performs Multiply and Square once
- Multiply and Square performs mostly MMM heavily in a loop
- Focus on optimizing MMM

Optimizations Performed

- v0 – MMM Operator Strength Reduction
 - Multiplication is not required in MMM
 - One of the operands is always 0 or 1
- v1 – Further MMM Optimization
 - Replacing right shifts by *i*-bit to only by 1 or 32
 - Storing only relevant 32-bit chunks of ***uint96_t***
- v2 – Optimizing ***uint96_t*** addition and subtraction
 - Loop Unrolling
- v3 - Rewriting ***rshift_uint96*** by *i*-bit function
 - To two functions handling only right shifts by 1 or 32
 - Remove if branches required to handle shifts ranging 0-96
- v4 - Adding Register Keywords

Original MMM

Optimizations performed

- Operator Strength Reduction

```
uint96_t mmm(uint96_t X, uint96_t Y, uint96_t M, uint32_t numBits)
{
    uint96_t T_low = {0};
    uint32_t T_high = 0;

    int m = numBits;

    for (int i = 0; i < m; i++)
    {
        uint32_t n = (T_low.value[2] & 1) ^ ((rshift_uint96(X, i).value[2] & 1) & (Y.value[2]
& 1));

        if (rshift_uint96(X, i).value[2] & 1){
            T_low = add_uint96(T_low, Y);

            if (cmp_uint96(T_low, Y)){
                T_high += 1;
            }
        }

        if (n){
            T_low = add_uint96(T_low, M);

            if (cmp_uint96(T_low, M)){
                T_high += 1;
            }
        }

        T_low = rshift_uint96(T_low, 1);
        T_low.value[0] += T_high << 31;
        T_high >>= 1;
    }

    if (T_high & 1 || !cmp_uint96(T_low, M)){
        T_low = sub_uint96(T_low, M);
    }

    return T_low;
}
```

Final MMM

Optimizations performed

- Operator Strength Reduction
- Storing relevant 32-bit sizes
- Replacing shifts by i -bits to 1 or 32
- Register Keyword

```
uint96_t mmm(uint96_t X, uint96_t Y, uint96_t M, int numBits)
{
    register uint96_t T_low = {0};
    register uint32_t T_high = 0;
    register uint32_t n;
    register uint32_t X_32local;
    register uint32_t Y_0 = Y.low & 1;

    for (int i = 0; i < numBits; i++)
    {
        if(!(i % 32)){
            X_32local = X.low;
            X = rshift32_uint96(X);
        }
        n = (T_low.low & 1) ^ ((X_32local & 1) & Y_0);

        if (X_32local & 1){
            T_low = add_uint96(T_low, Y);

            if (cmp_uint96(T_low, Y)){
                T_high += 1;
            }
        }

        if (n){
            T_low = add_uint96(T_low, M);

            if (cmp_uint96(T_low, M)){
                T_high += 1;
            }
        }

        T_low = rshift1_uint96(T_low);
        T_low.high += T_high << 31;
        T_high >>= 1;

        X_32local >>= 1;
    }
    if (T_high & 1 || !cmp_uint96(T_low, M)){
        T_low = sub_uint96(T_low, M);
    }
    return T_low;
}
```


MMM Optimization – Replace shifts by i to by 1 or 32

- Compiler performs in-lining for ***rshift_uint96***
- Before: ***rshift_uint96*** by loop counter i ranging from 0 to 96
 - 271653 calls/branches to ***rshift_uint96*** in assembly
- After: ***rshift_uint96*** by constants 1 and 32
 - no calls/branches to ***rshift_uint96*** in assembly

MMM Optimization – Assembly changes

```
283     mov     r0, r7
284     bl      rshift_uint96
285     ldm     r7, {r0, r1, r2}
286     stm     r4, {r0, r1, r2}
287     mov     r9, r2
288     add     r6, r0, fp, lsl #31
289     lsr     fp, fp, #1
290     add     r10, r10, #1
291     ldr     r3, [sp, #20]
292     cmp     r3, r10
293     beq     .L45
294 .L34:
295     str     r10, [sp]
296     ldm     r8, {r1, r2, r3}
297     add     r0, sp, #56
298     bl      rshift_uint96
299     ldr     r5, [sp, #64]
300     ldr     r3, [sp, #16]
301     and     r5, r5, r3
302     eor     r5, r5, r9
303     and     r5, r5, #1
304     str     r10, [sp]
305     ldm     r8, {r1, r2, r3}
306     add     r0, sp, #68
307     bl      rshift_uint96
308     ldr     r3, [sp, #76]
309     tst     r3, #1
310     beq     .L28
311     str     r6, [r4]
```

```
282 .L40:
283     add     r7, r7, #1
284 .L32:
285     lsl     r0, r2, #31
286     add     r1, r0, r1, lsr #1
287     lsl     r0, r3, #31
288     add     r2, r0, r2, lsr #1
289     lsr     r3, r3, #1
290     add     r3, r3, r7, lsl #31
291     lsr     r7, r7, #1
292     lsr     r6, r6, #1
293     add     r8, r8, #1
294     cmp     r9, r8
295     beq     .L47
296 .L35:
297     tst     r8, #31
298     beq     .L48
299 .L28:
300     and     r4, r6, r10
301     eor     r4, r4, r1
302     and     r4, r4, #1
303     tst     r6, #1
304     beq     .L29
305     str     r3, [sp, #68]
306     str     r2, [sp, #72]
307     str     r1, [sp, #76]
308     str     fp, [sp, #128]
309     add     r3, sp, #120
310     ldm     r3, {r0, r1, r2}
```

uint96_t Optimization – Loop unrolling

Before loop unrolling to 96-bit addition and subtraction

- Internal array indexing in loop prevents register usage
- No In-lining to ***add_uint96*** and ***sub_uint96***
 - 89433 calls to ***add_uint96***
 - 191 calls to ***sub_uint96***

After loop unrolling to 96-bit addition and subtraction

- Register keyword can be used
- In-lining by compiler to ***add_uint96*** and ***sub_uint96***
 - 0 calls to both operations

Optimizations that Reduced Performance

- v2neon - Neon Addition
 - 96-bit addition body was replaced with Neon addition
 - Costly overhead of storing and loading into Neon data types/registers
- v5 - Loop-unrolling for MMM
 - Unrolling of five times to target 95-bit keys
 - Reduced performance by 33.75%
 - Tight dependencies = No parallelization
 - Larger code = More instruction cache misses

Cycle Count Averages (Rough Estimation)

| Version | Optimizations Performed | Cycle Count Averages (Rough Est.) | | |
|---------------|---|-----------------------------------|---------|--------|
| | | -O0 | -O1 | -O3 |
| v0 | Operator Strength Reduction <ul style="list-style-type: none">Avoid 96-bit multiplication as one operand is always 0 or 1 | 43017.8 | 27207.2 | 6116.6 |
| v1 | + MMM Optimization <ul style="list-style-type: none">Storing only relevant 32-bit chunks of uint96_trshift_uint96 in-lining by replacing shifts by <i>i</i>-bit to 1 or 32 | 31139.2 | 12256.2 | 4931.2 |
| v2 | + uint96_t Addition and Subtraction Optimization <ul style="list-style-type: none">Loop UnrollingAvoided internal array indexing | 29255.0 | 4022.2 | 5023.0 |
| v2neon | + Neon Addition (Not pursued / Removed for next version) | 33277.4 | 13961.6 | 8141.6 |
| v3 | + Rewriting rshift_uint96 to two functions handling only right shifts by 1 or 32 | 27711.8 | 3925.6 | 5251.8 |
| v4 | + Register Keyword | 21813.4 | 4254.2 | 4830.0 |

Hardware-assisted Solution Analysis

- A large portion of time is spent on ***rshift_uint96*** and ***add_uint96***
- ***rshift_uint96*** costs only approximately 5-8 instructions
- ***add_uint96*** costs around 18 instructions
 - Operands however are two 96-bit integers
 - Would require multiple hardware
- Larger chunks of operations that uses these 96-bit operations considered
 - Tight dependencies
 - Difficult put larger chunk of code into functions (especially of 32-bit operands)

Conclusion

- Important to understand software optimization techniques
 - Benefited strongest level of safe optimization
- Test runtime after each optimization
 - Beneficial in theory, but difficult to predict cache misses or how compiler interprets code
- -O1 flag performed better than the -O3 flag after manual optimizations
 - More optimization flags does not always mean faster program
 - Manual software optimizations can result in better performance

Any Questions?